

CS-583: Deep Learning

BERT Model

Abdul Rafae Khan

Department of Computer Science
Stevens Institute of Technology
akhan4@stevens.edu

November 2, 2023

What is BERT?

- BERT [1] is for **pre-training** Transformer's [2] encoder.
- How?
- Predict masked word.
- Predict next sentence.

Reference


1. Devlin, Chang, Lee, and Toutanova. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *ACL*, 2019.
2. Vaswani and others. [Attention is all you need](#). In *NIPS*, 2017.

Task 1: Predict Masked Words

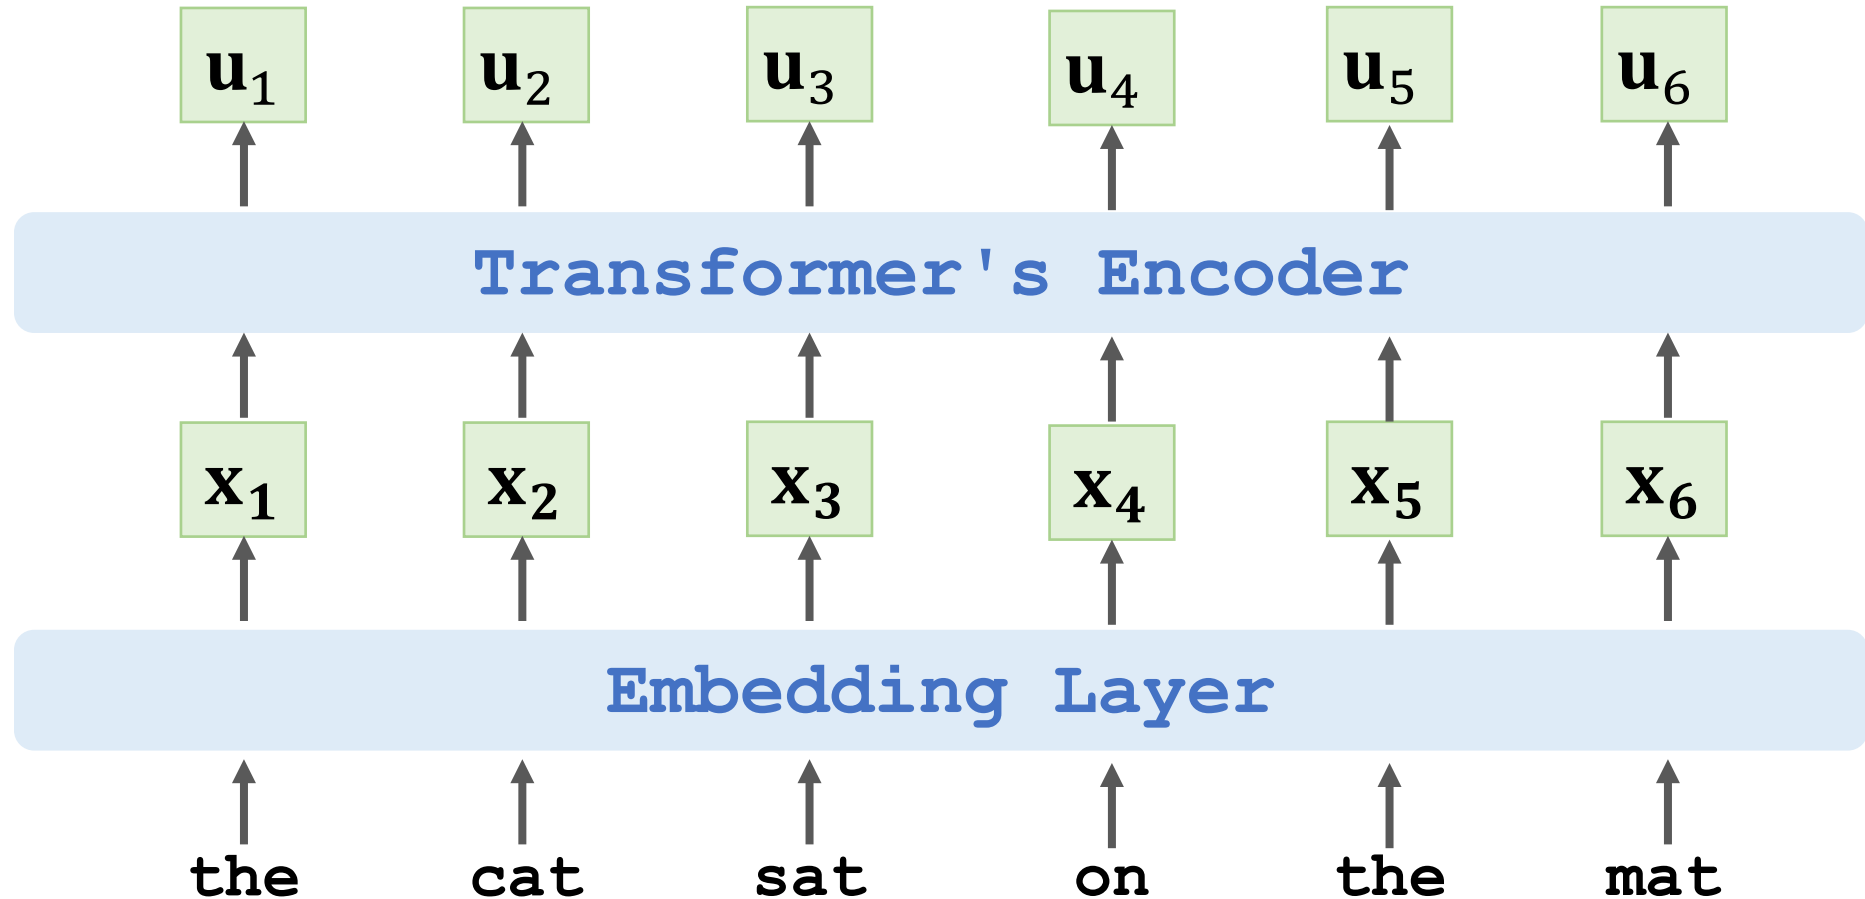
Revisit Transformer's Encoder

the cat sat on the mat

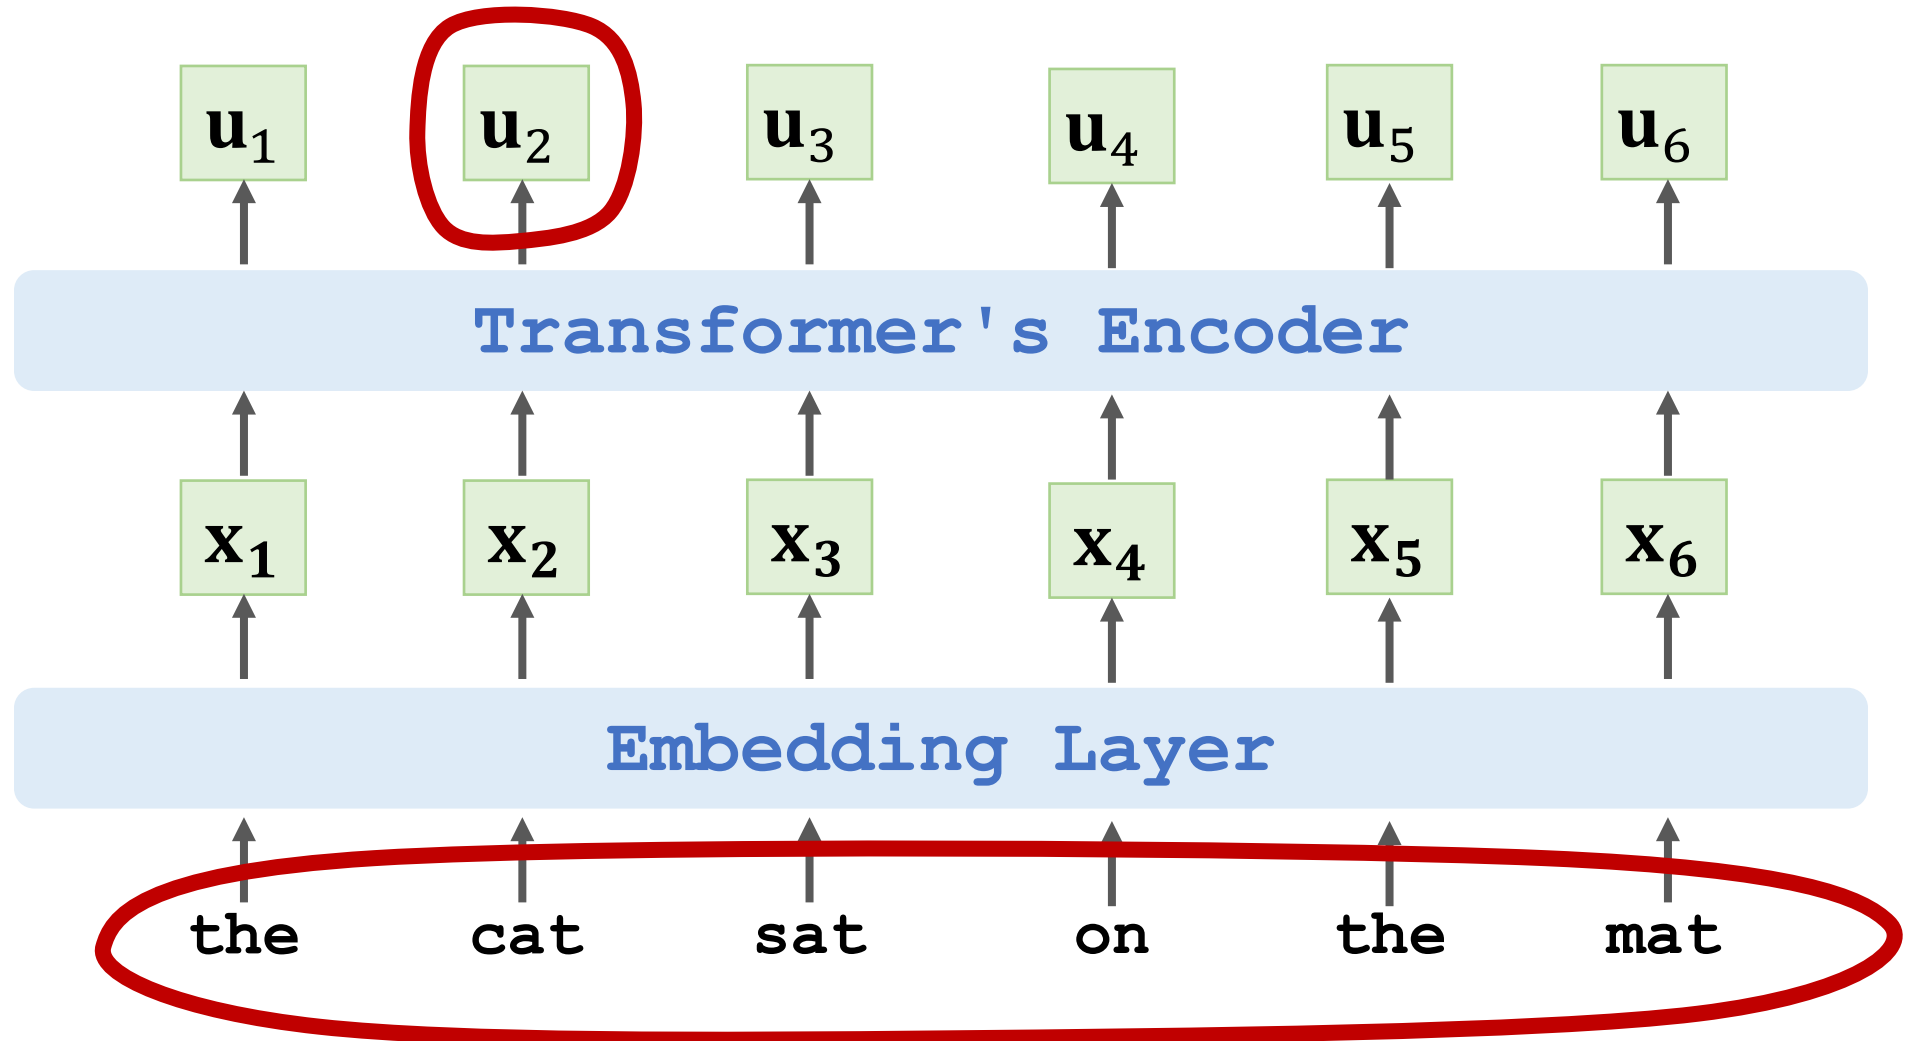
Revisit Transformer's Encoder



Revisit Transformer's Encoder



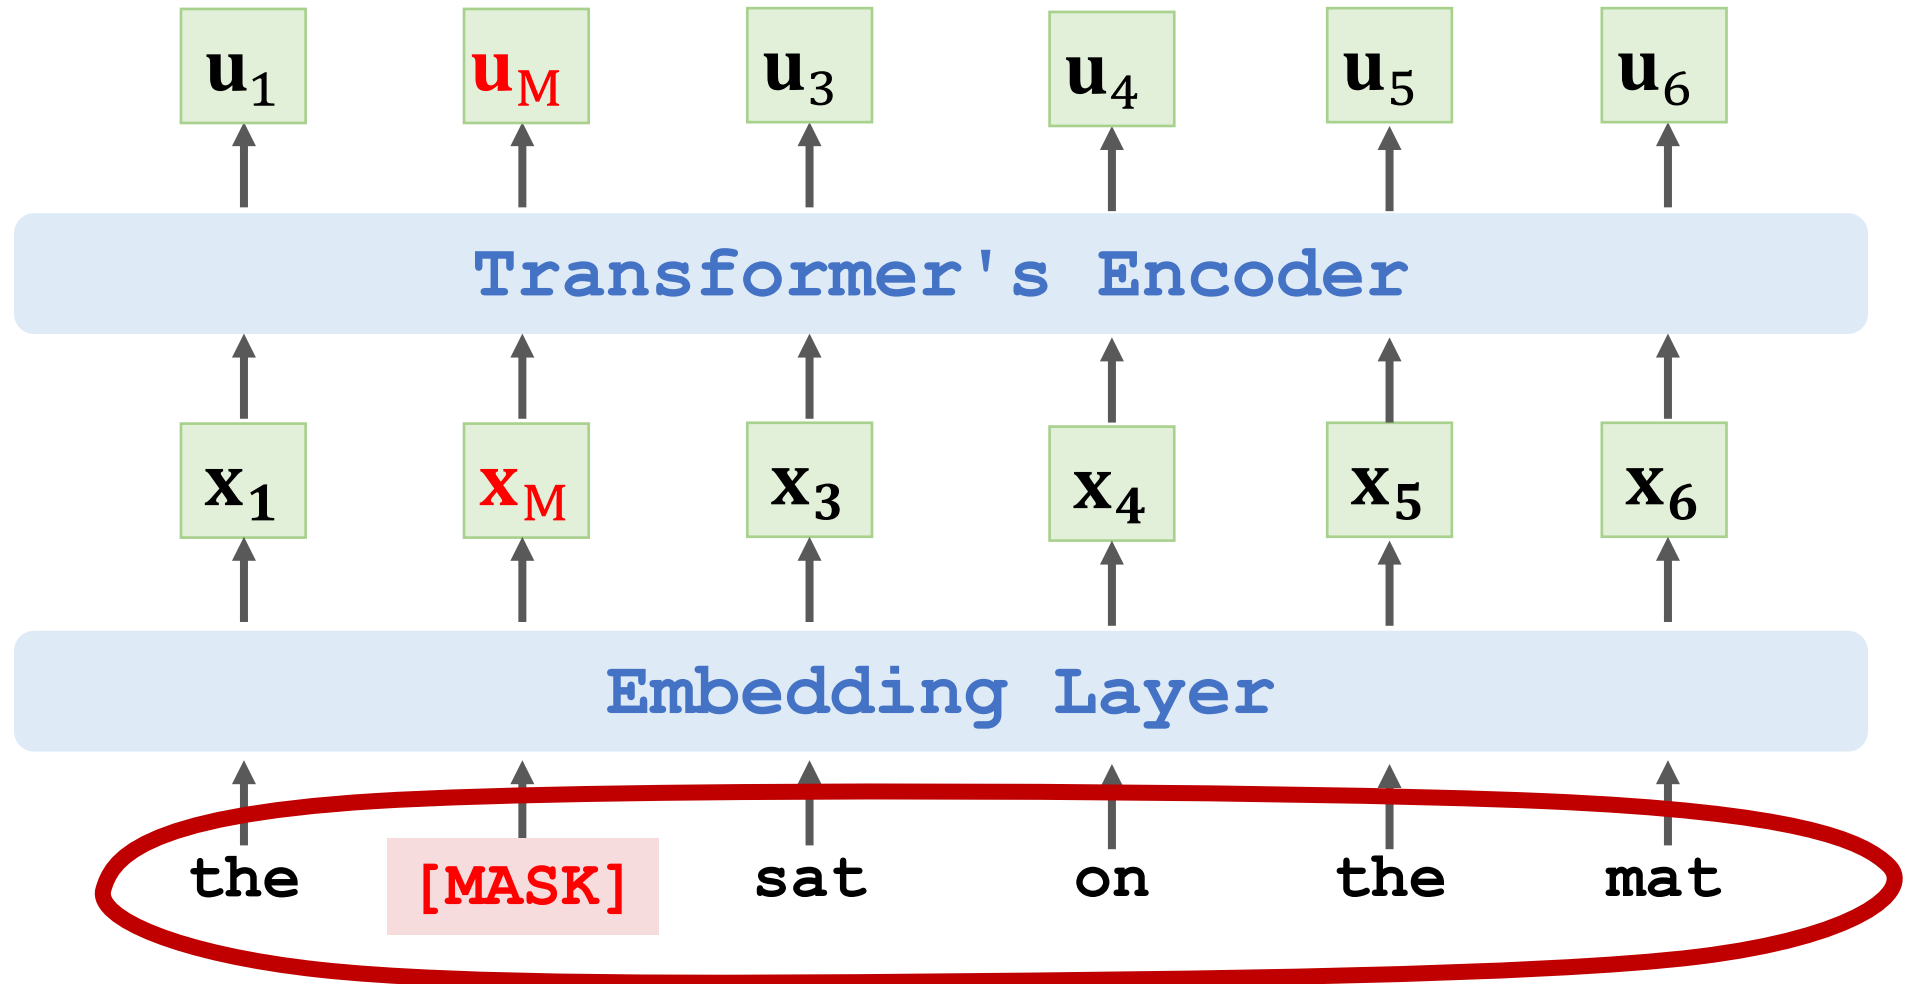
Revisit Transformer's Encoder

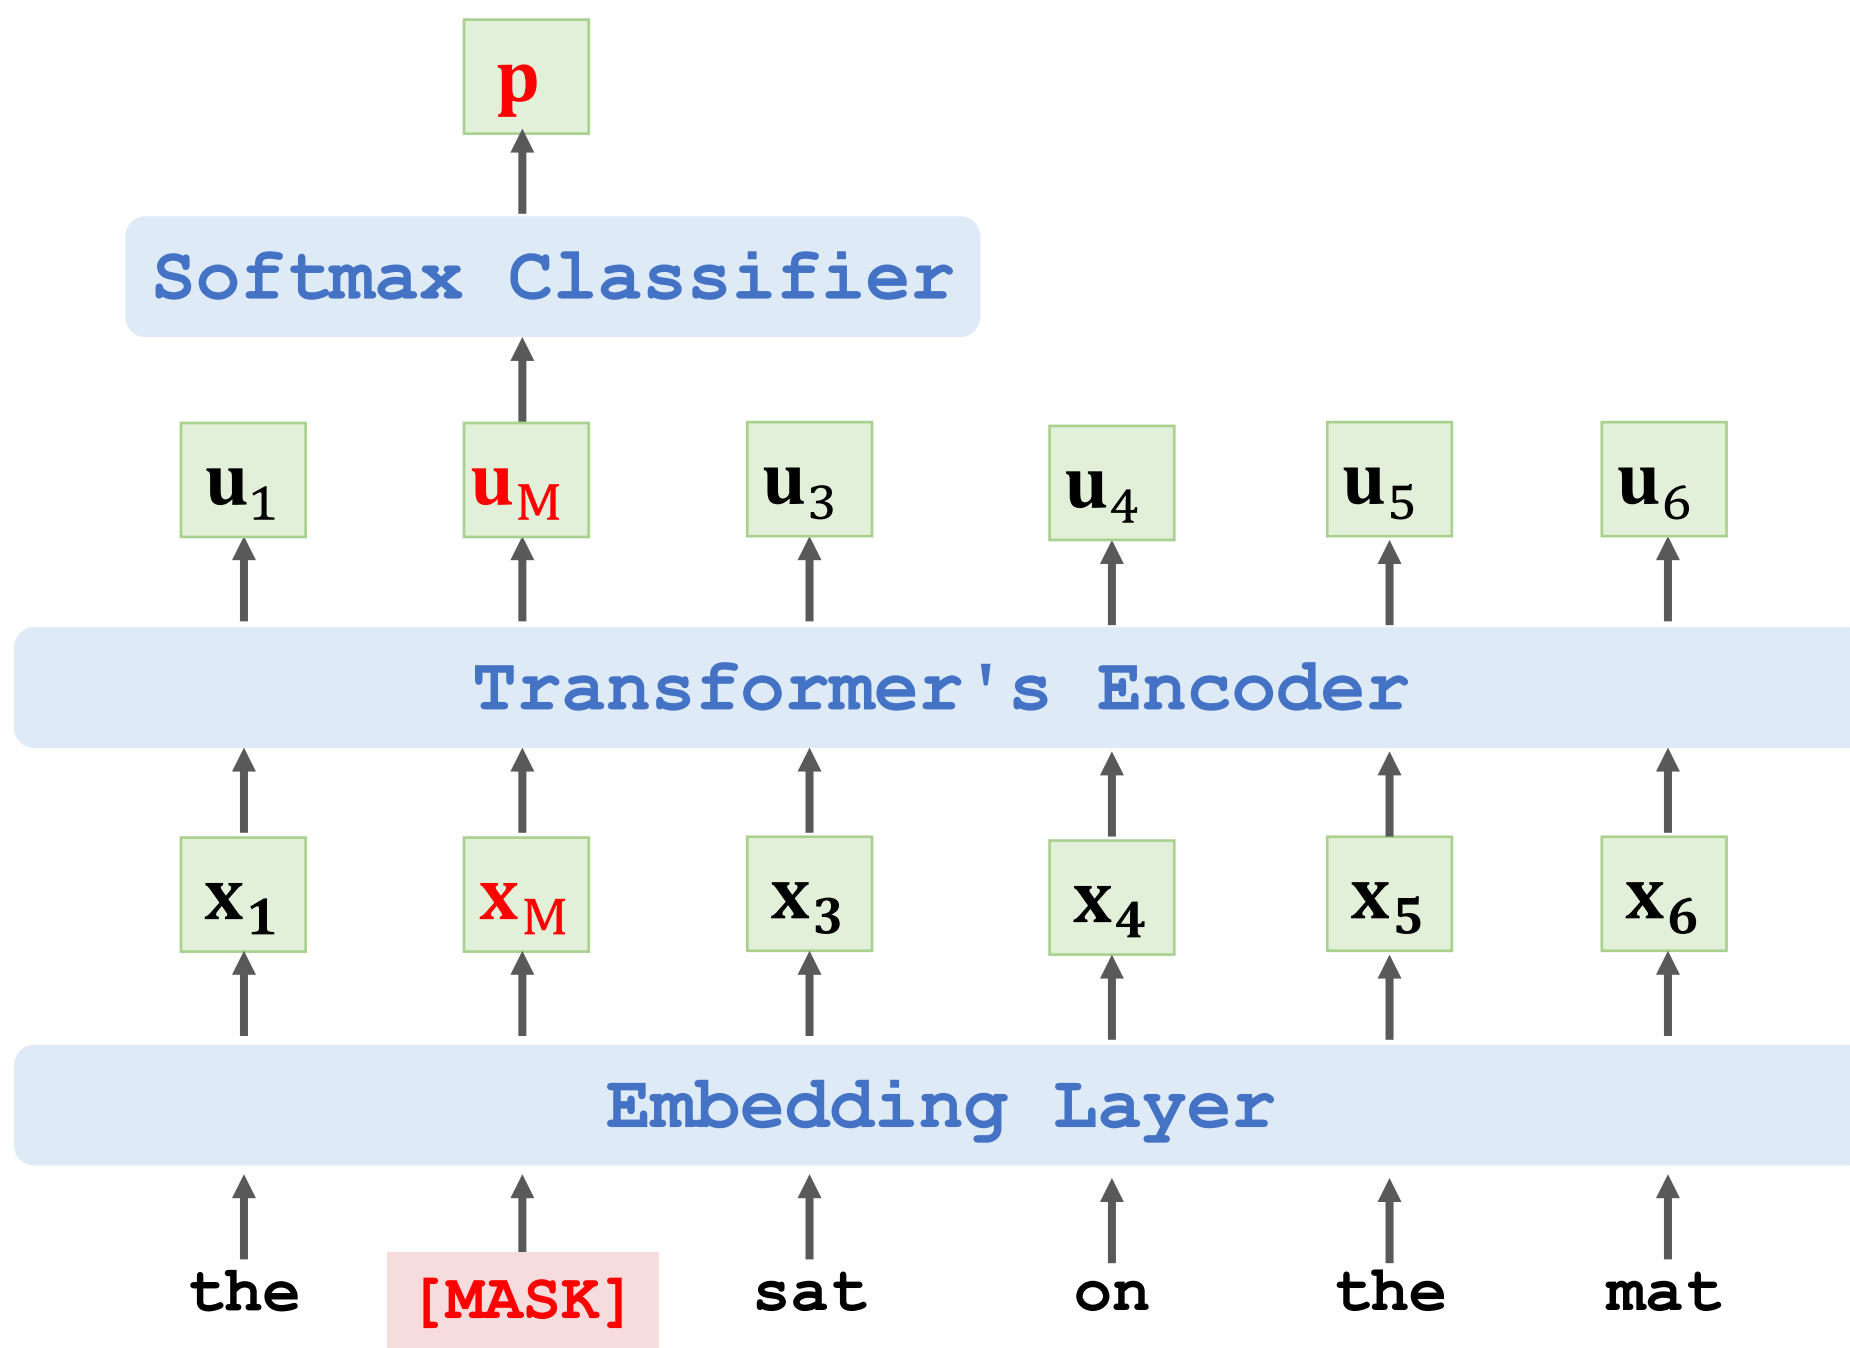


Randomly mask a word

- “The _____ sat on the mat”
- What is the masked word?

Randomly mask a word





Predict the masked word

- **e**: one-hot vector of the masked word “cat”.
- **p**: output probability distribution.
- $\text{Loss} = \text{CrossEntropy}(\mathbf{e}, \mathbf{p})$.
- Performing one gradient descent to update the model parameters.

Task 2: Predict the Next Sentence

Predict the next sentence

- Given the sentence:

"calculus is a branch of math".

- Is this the next sentence?

"it was developed by newton and leibniz"

Predict the next sentence

- Given the sentence:

"calculus is a branch of math".

- Is this the next sentence?

"it was developed by newton and leibniz"

- Is this the next sentence?

"panda is native to south central china"

Input Representation

- **Input:**

[CLS] "calculus is a branch of math"

[SEP] "it was developed by newton and leibniz"

Input Representation

- **Input:**

[CLS] "calculus is a branch of math"

[SEP] "it was developed by newton and leibniz"

- [CLS] is a token for classification.

- [SEP] is for separating sentences.

Input Representation

- **Input:**

[CLS] "calculus is a branch of math"

[SEP] "it was developed by newton and leibniz"

Target: true

Input Representation

- **Input:**

[CLS] "calculus is a branch of math"

[SEP] "panda is native to south central china"

- **Target:** false

Predict the next sentence

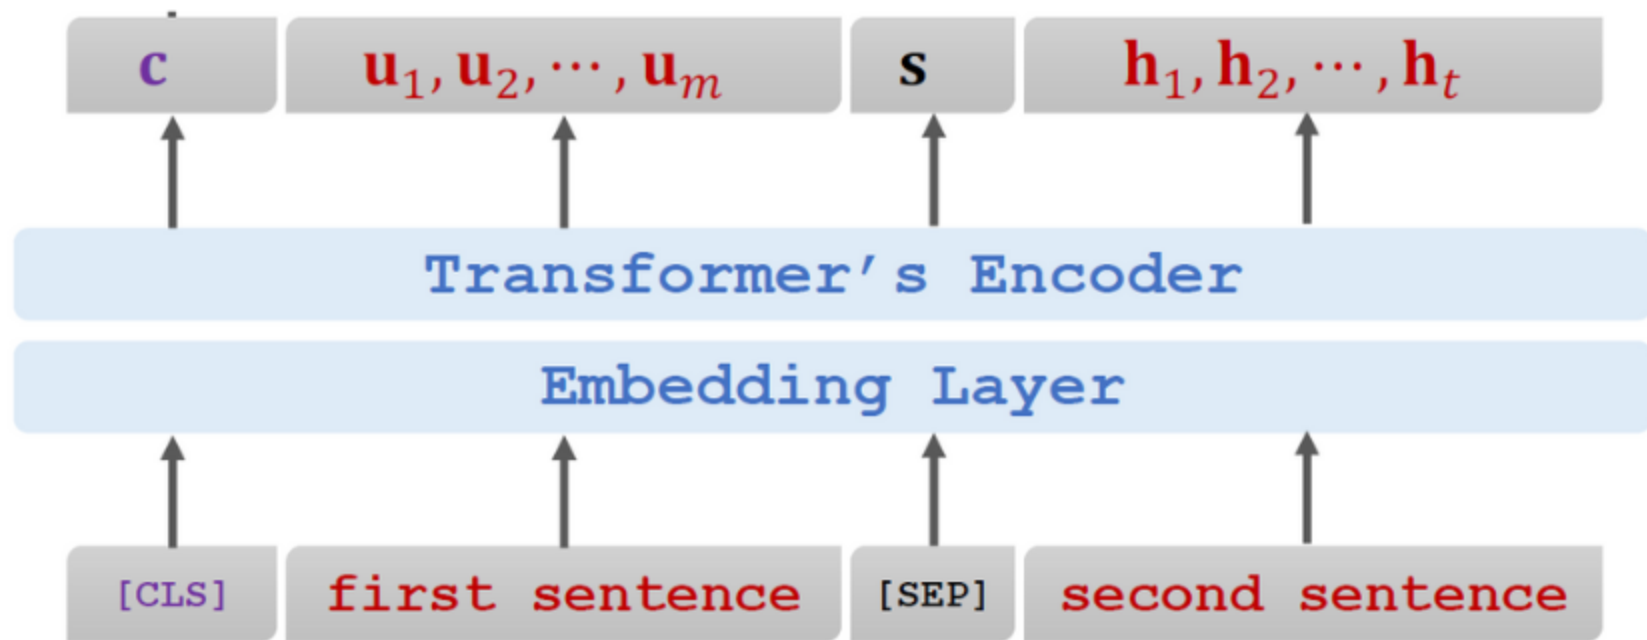
[CLS]

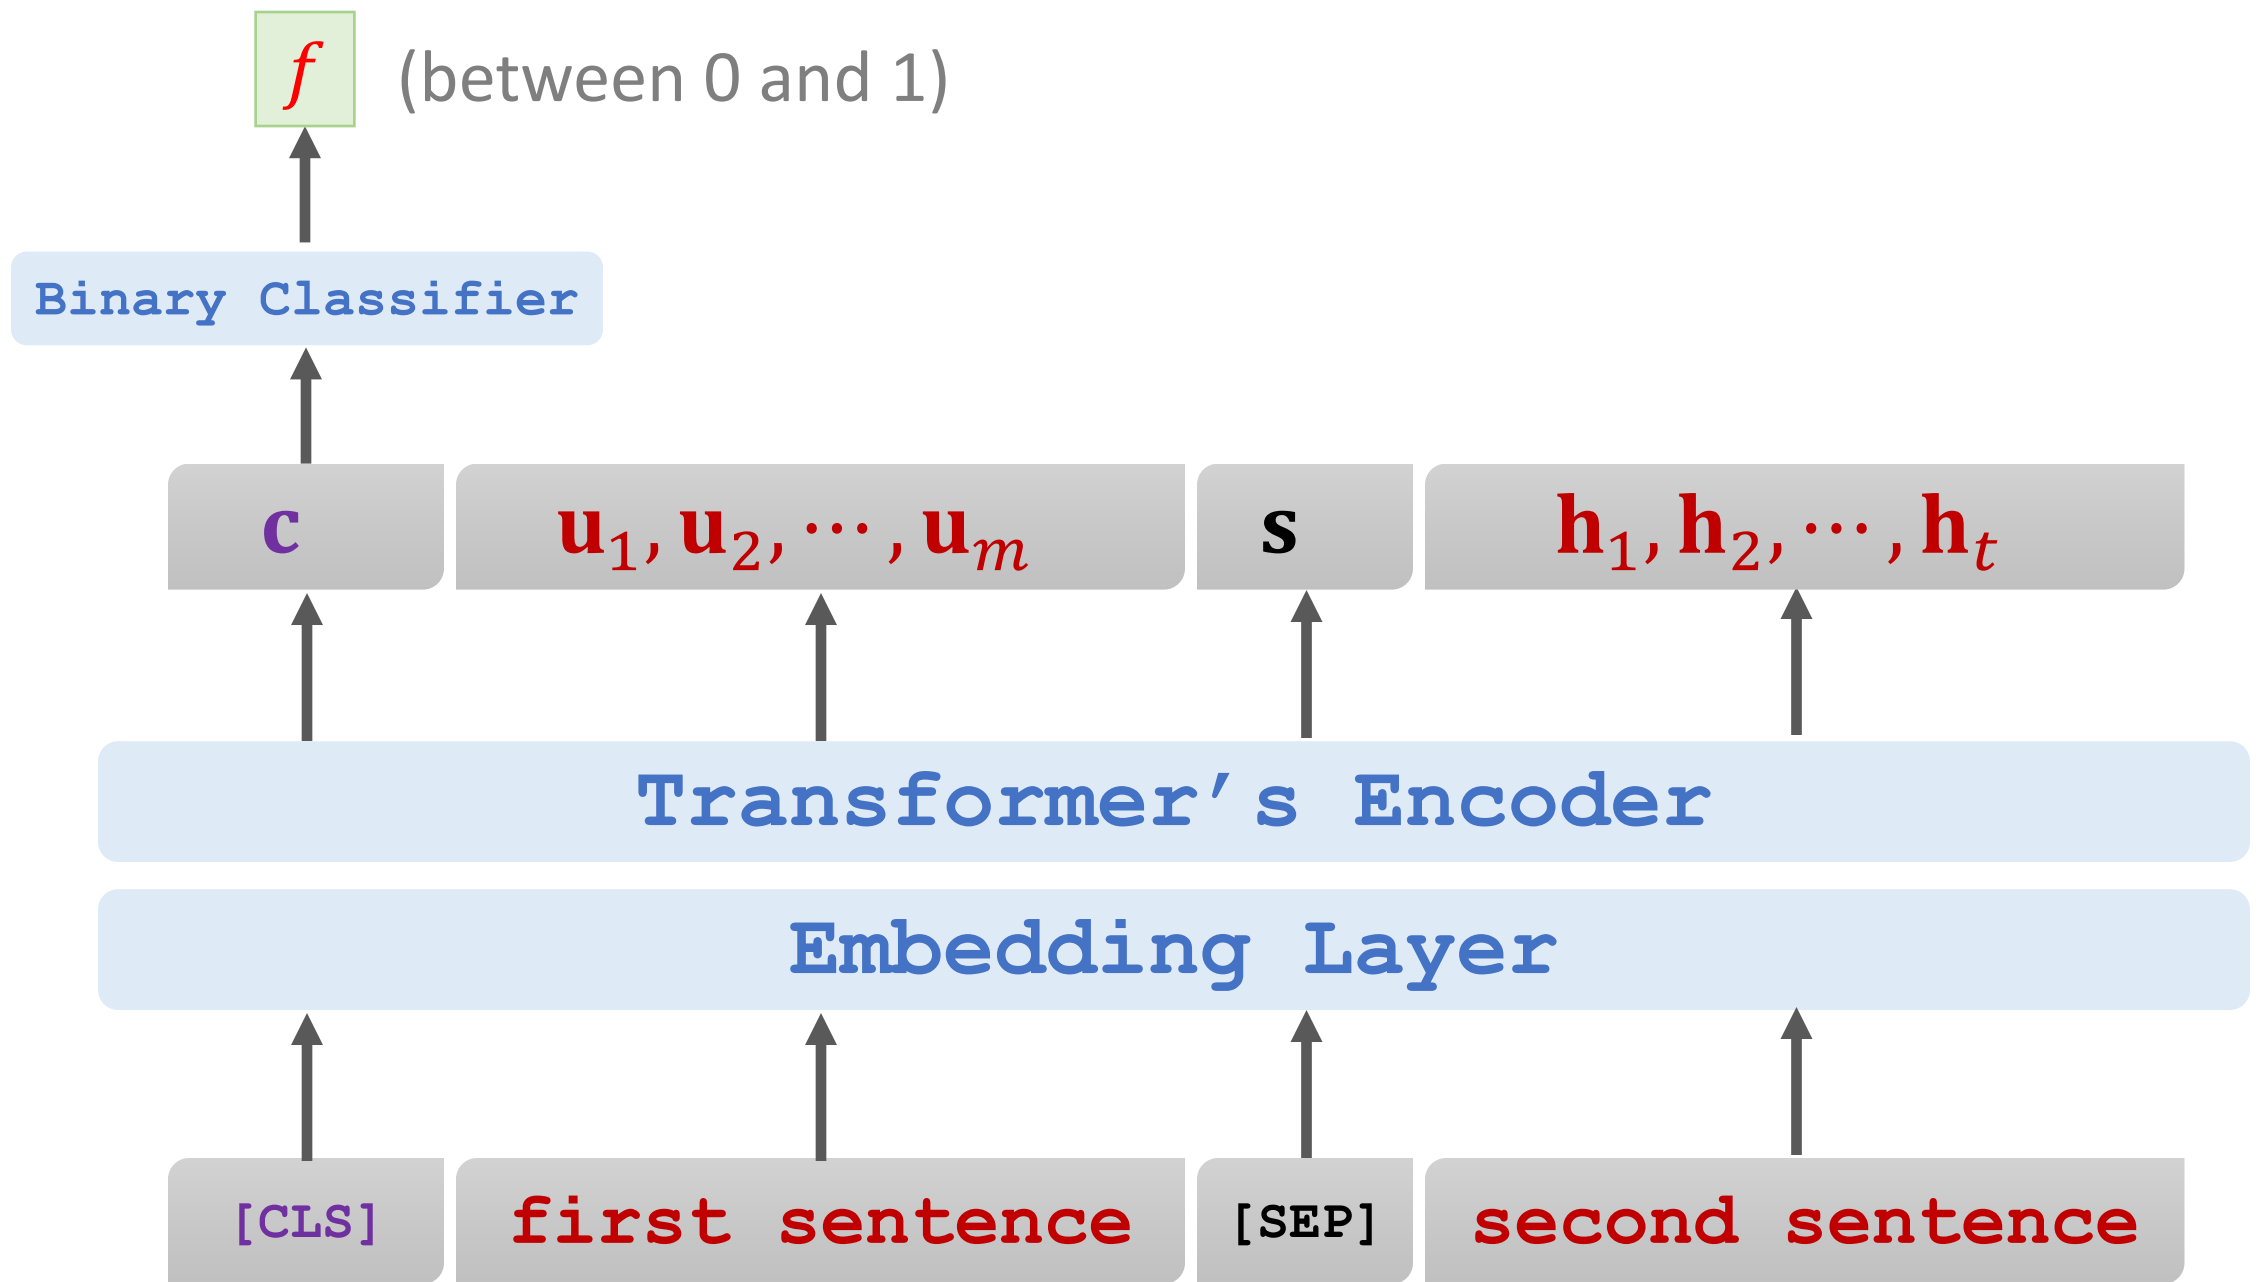
first sentence

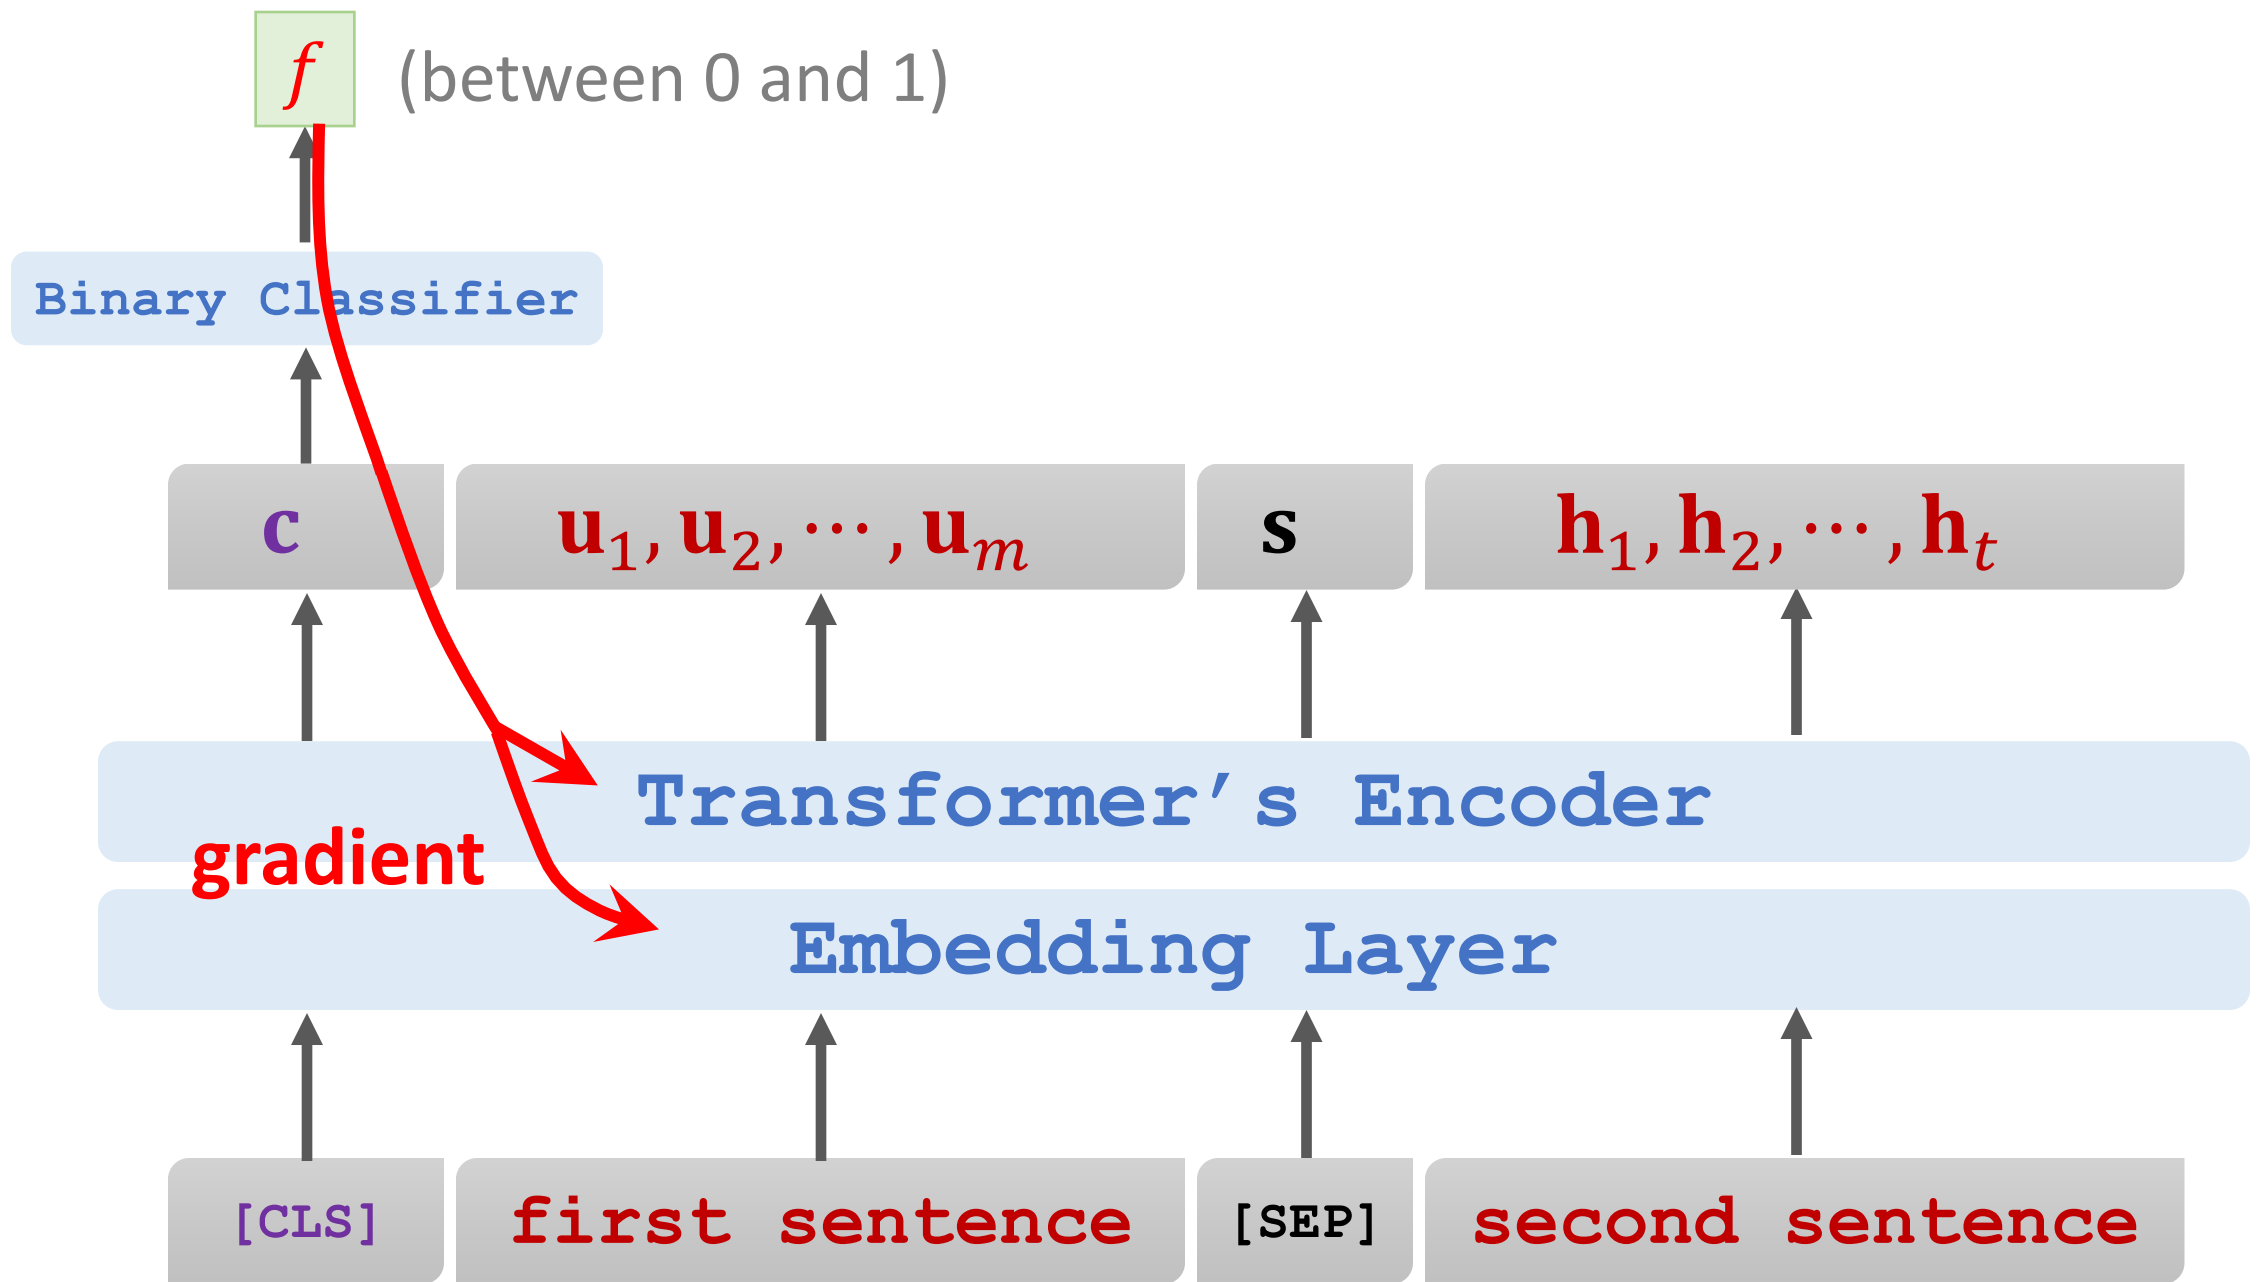
[SEP]

second sentence

Predict the next sentence







Why is the classification task helpful?

- The task helps improve the embedding layer.
- The task helps improve the self-attention layers.

Combining the two methods

Input Representation

- Input:

"[CLS] calculus is a [MASK] of math
[SEP] it [MASK] developed by newton and leibniz"

Input Representation

- **Input:**

"[CLS] calculus is a [MASK] of math
[SEP] it [MASK] developed by newton and leibniz".

- **Targets:**

true, "branch" "was"

Input Representation

- **Input:**

"[CLS] calculus is a branch of math
[SEP] panda is native to [MASK] central china".

- **Targets:** false, "south".

Training

- ➡ • **Loss 1** is for binary classification (i.e., deciding if the next sentence is real.)
- ➡ • **Loss 2** and **Loss 3** are for multi-class classification (i.e., predicting the masked words.)
- **Objective function** is the **sum** of the three **loss functions**.
- Update model parameters by performing one gradient descent.

Data

- BERT does not need manually labeled data. (Nice! Manual labeling is expensive.)
- Use large-scale data, e.g., English Wikipedia (2.5 billion words.)
- Randomly mask 15% words (with some tricks.)
- 50% of the next sentences are real. (The other 50% are fake.)

Cost of Computation



- BERT Base

- 110M parameters.
- 16 TPUs, 4 days of training (without hyper-parameter tuning.)



- BERT Large

- 235M parameters.
- 64 TPUs, 4 days of training (without hyper-parameter tuning.)

The End