

# CS-583: Deep Learning Reinforcement Learning


Abdul Rafae Khan

Department of Computer Science  
Stevens Institute of Technology  
*akhan4@stevens.edu*

November 9, 2023

# Search Problem


---



# Deterministic Actions

---

Actions from a give state are deterministic  
 $Succ(s, a)$  is always the same state  $s'$




# Stochastic Actions

---




Actions from a give state are probabilistic (stochastic)

$Succ(s, a, t)$  denotes the next state given the current state  $s$  and action  $a$  taken at the time  $t_i$

It can either be state  $B$  with a probability  $p_i$  or state  $C$  with probability  $p_j$



# Applications



# Markov Decision Process

---

## Game:


The player starts with \$0 as the prize money.

In each round, the player can take two steps:

- Quit and take \$10
- Answer a question
  - Correctly answer with a probability of  $\frac{2}{3}$ , get \$4 prize and move to the next round
  - Otherwise get \$4 prize and end the game

# Markov Decision Process

---



# Markov Decision Process

---

## Gameshow:

The player starts with 0 as the prize money.

In each round, the player can take two steps:

- Quit and take \$10
- Answer a question
  - Correctly answer with a probability of  $\frac{2}{3}$  and move to the next round
  - Otherwise take \$4 and end the game


**What is the best strategy for the game?**



# Markov Decision Process

---


If our policy is to 'answer':



# Markov Decision Process

---

If our policy is to 'answer':




Expected Utility:

$$\frac{1}{3}(4) + \frac{2}{3} \cdot \frac{1}{3}(8) + \frac{2}{3} \cdot \frac{2}{3} \cdot \frac{1}{3}(12) + \dots = 12$$

# Markov Decision Process

---

If our policy is to 'quit':



Expected Utility:

$$1(10) = 10$$

# Search problem

---

$s_{start}$ : start state

**Actions( $s$ )**: all possible actions from state  $s$

**Succ( $s, a$ )**: next possible states given action  $a$  is taken from state  $s$

**Cost( $s, a$ )**: cost of transition from state  $s$  by taking action  $a$

**IsEnd( $s$ )**: is  $s$  a goal state

# Markov Decision Process

---

$s_{start}$ : start state

**Actions**( $s$ ): all possible actions from state  $s$

**T**( $s, a, s'$ ): probability of  $s'$  if action  $a$  is taken from state  $s$

**Reward**( $s, a, s'$ ): reward from the transition  $s$  to  $s'$

**IsEnd**( $s$ ): is  $s$  a goal state

$0 \leq \gamma \leq 1$ : discount factor (default: 1)

# Markov Decision Process

---

Total transition probability:  $\sum_{s'} T(s, a, s') = 1$

Discount factor  $\gamma$  is based on how much we value the future reward

# Markov Decision Process

---

$Succ(s, a) \rightarrow T(s, a, s')$

$Succ(s, a)$  can be considered as a special case of transition probability

$$T(s, a, s') = \begin{cases} 1 & \text{if } s' = Succ(s, a) \\ 0 & \text{otherwise} \end{cases}$$

# Markov Decision Process

---

$\text{Cost}(s, a) \rightarrow \text{Reward}(s, a, s')$

Instead of minimizing the cost, we maximize the reward

Negating one is equivalent to the other



# Markov Decision Process

---

$T(s, a, s')$ : probability of  $s'$  if action  $a$  is taken from state  $s$

$s$	$a$	$s'$	$T(s, a, s')$
$s_{start}$	Quit	$s_{end}$	1
$s_{start}$	Question	$s_{end}$	$1/3$
$s_{start}$	Question	$s_{start}$	$2/3$

# Markov Decision Process

---

$T(s, a, s')$ : probability of  $s'$  if action  $a$  is taken from state  $s$

$s$	$a$	$s'$	$T(s, a, s')$
$s_{start}$	Quit	$s_{end}$	1
$s_{start}$	Question	$s_{end}$	$1/3$
$s_{start}$	Question	$s_{start}$	$2/3$

To re-iterate:

Sum of probabilities from a given state  $s$  by making an action  $a$  is 1

$$\sum_{s' \in \text{states}} T(s, a, s') = 1$$

Successors: states  $s'$  where  $T(s, a, s') > 0$

# Markov Decision Process

---

$T(s, a, s')$ : probability of  $s'$  if action  $a$  is taken from state  $s$

$s$	$a$	$s'$	$T(s, a, s')$
$s_{start}$	<i>Quit</i>	$s_{end}$	1
$s_{start}$	<i>Question</i>	$s_{end}$	$1/3$
$s_{start}$	<i>Question</i>	$s_{start}$	$2/3$

Sum of probabilities from a given state  $s$  by making an action  $a$  is 1

# Policy

---

**Policy:** gives an action  $a$  for a given  $\pi : s \rightarrow a$

For deterministic search problems, we wanted the optimal sequence of actions from start to goal

For MDP, we want the optimal policy  $\pi^* : s \rightarrow a$  which maximizes the reward

$\text{Reward}(s, a, s')$

# Grid World!


---

Our world is  $3 \times 4$  grid

Start state is at (0,0)

Reward +1 at (4,3)

Reward -1 at (4,2)




# Grid World!

---

For any state, three possible moves

- up: 0.8
- left: 0.1
- right: 0.1



# Grid World!

---

3	→	→	→	+1
2	↑		↑	-1
1	↑	←	↖	←
	1	2	3	4

Optimal policy for  $\gamma < -0.04$

There are two optimal policies for state (3,1)

# Discount

---

## Additive discount utility

Let say the path is  $s_0, a_1 r_1 s_1, a_2 r_2 s_2, \cdot$  (sequence of state, action, and reward)

The utility with discount  $\gamma$  is:

$R(s, a, s') + \gamma R(s, a, s') + \gamma^2 R(s, a, s') + \dots$  where  $\gamma \in [0, 1]$

$\gamma$  is based on how important current reward is compared to the future reward



# Discount

---

Solving the problem of infinite stream of rewards

Geometric series:  $1 + \gamma + \gamma^2 + \dots = 1/(1 - \gamma)$

Assume rewards bounded by  $\pm R_{max}$

Then  $r_0 + \gamma_1 r_1 + \gamma_2 r_2 + \dots$  is bounded by  $\pm R_{max}/(1 - \gamma)$

# Policy Evaluation

---

The **utility** is the discounted sum of rewards on the path.

Optimal policy:  $\pi^*(s)$  = optimal actions from state  $s$

It gives highest  $U_\pi(s)$  for any  $\pi$

$$U_\pi(s) = R(s, a, s') + \gamma R(s, a, s') + \gamma^2 R(s, a, s') + \dots$$

# Policy Evaluation


---

For a given policy  $\pi$ , we have two variable associated with it:

- Value of the policy  $V_{\pi}(s)$
- Q-value of the policy  $Q_{\pi}(s, \pi(s))$

# Markov Decision Process

---



# Policy Evaluation

---

For a given policy  $\pi$ , we have two variable associated with it:

- Value of the policy  $V_{\pi}(s)$
- Q-value of the policy  $Q_{\pi}(s, \pi(s))$

The value can be thought of as the label for the nodes representing the states and the Q-value as the label for the chance nodes

# Policy Evaluation

---

**Value** is the expected utility from following policy  $\pi$  from state  $s$

**Q-value** is the expected utility of taking action  $a$  from state  $s$ , and then following policy  $\pi$ .

$$V_{\pi}(s) = E[V_{\pi}(s)] = \begin{cases} 0 & \text{if } isEnd(s) \\ Q_{\pi}(s) & \text{otherwise} \end{cases}$$

$$Q_{\pi}(s) = \sum_{s'} T(s'|s, a) [R(s, a, s') + \gamma V(s')]$$

# Policy Evaluation

Let the policy  $\pi$  be 'Answer':

$$V_{\pi}(s_{end}) = 0$$

$$V_{\pi}(s_{start}) = Q_{\pi}(s_{start}, \text{Answer})$$


$$= \frac{1}{3}(4 + V_{\pi}(s_{end})) + \frac{2}{3}(4 + V_{\pi}(s_{start}))$$

$$\Rightarrow V_{\pi}(s_{start}) = \frac{1}{3}(4) + \frac{2}{3}(4 + V_{\pi}(s_{start}))$$

Closed form solution:

$$3V_{\pi}(s_{start}) = 4 + 2 \cdot 4 + 2V_{\pi}(s_{start})$$

$$V_{\pi}(s_{start}) = 12$$



# Policy Evaluation

---

Given the recursion  $V^*(s) = \max_a Q^*(s, a)$

**Value:**

$$V^*(s) = \max_{a \in \text{Actions}(s)} \sum_{s'} \{P(s'|s, a)[R(s, a, s') + \gamma V(s')]\}$$

**Q-value:**

$$\begin{aligned} Q^*(s, a) &= \sum_{s'} \{P(s'|s, a)[R(s, a, s') + \gamma V(s')]\} \\ &= \sum_{s'} \{P(s'|s, a)[R(s, a, s') + \gamma \max_{a'} Q(s', a')]\} \end{aligned}$$



# Markov Decision Process

---

Solving MDPs:

- Value Iteration
- Policy Iteration

# Policy Iteration

---

$$V_{\pi}^{(0)}(s) \leftarrow 0$$

for  $i = 1 \cdots t_{max}$

for each state  $s$

$$V_{\pi}^{(t)}(s) \leftarrow \sum_{s'} T(s'|s, a) [R(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]$$

# Policy Iteration

---

$V_{\pi}^{(0)}(s) \leftarrow 0$   
for  $i = 1 \cdots t_{max}$   
  for each state  $s$   
     $V_{\pi}^{(t)}(s) \leftarrow \underbrace{\sum_{s'} T(s'|s, a) [R(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]}_{Q_{\pi}^{(t-1)}(s)}$

# Policy Evaluation

---

How many iterations ( $t_{max}$ )?

Repeat until there is no/very little change

$$\max_{s \in \text{states}} |V_{\pi}^{(t)}(s) - V_{\pi}^{(t-1)}(s)| \leq \epsilon$$

Only save the last two iterations,  $V_{\pi}^{(t)}$  &  $V_{\pi}^{(t-1)}$

# Policy Iteration

---

```
 $V_{\pi}^{(0)}(s) \leftarrow 0$   
for  $i = 1 \cdots t_{max}$   
  for each state  $s$   
     $V_{\pi}^{(t)}(s) \leftarrow \sum_{s'} T(s'|s, a)[R(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]$ 
```

**Total states:**  $S$

**Actions per state:**  $A$

**Total successor** (with  $T(s'|s, a) > 0$ ):  $S'$

**Complexity:**  $O(SS't_{max})$

# Policy Iteration

---

Let the policy  $\pi$  be 'Answer':

$$V_{\pi}^{(t)}(s_{end}) = 0$$

$$V_{\pi}^{(t)}(s_{start}) = \frac{1}{3}(4 + V_{\pi}^{(t-1)}(s_{end})) + \frac{2}{3}(4 + V_{\pi}^{(t-1)}(s_{start}))$$

Iteration ( $t$ )	$V_{\pi}^{(t)}(s_{end})$	$V_{\pi}^{(t)}(s_{start})$
0	0.00	0.00
1	0.00	4.00
2	0.00	6.67
3	0.00	8.44
100	0.00	12.00

$$V_{\pi}^{(t)}(s_{start}) = 12$$

# Optimal Value

---

Goal: try to get directly at maximum expected utility

$V_{opt}(s)$  = is the maximum value obtained by any policy

# Optimal Value

---

Given the recursion  $V_{opt}(s) = \max_a Q_{opt}(s, a)$

**Value:**

$$V_{opt}(s) = \max_{a \in \text{Actions}(s)} \sum_{s'} \{ T(s'|s, a) [R(s, a, s') + \gamma V_{opt}(s')] \}$$


**Q-value:**

$$\begin{aligned} Q_{opt}(s, a) &= \sum_{s'} \{ T(s'|s, a) [R(s, a, s') + \gamma V_{opt}(s')] \} \\ &= \sum_{s'} \{ T(s'|s, a) [R(s, a, s') + \gamma \max_{a'} Q_{opt}(s', a')] \} \end{aligned}$$



# Optimal Value

---



# Optimal Value

---

Policy evaluation used the action from a fixed policy  $\pi$

Now we pick the action which maximizes the Q-value  $Q_{opt}(s)$

$$V_{opt}(s) = \begin{cases} 0 & \text{if } isEnd(s) \\ \max_{a \in Actions(s)} Q_{opt}(s) & \text{otherwise} \end{cases}$$

$$Q_{opt}(s) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{opt}(s')]$$

# Optimal Policy

---

As for any state  $s$ ,  $Q_{\pi}(s)$  gives you the value of taking the policy  $\pi(s)$

Therefore, **Optimal policy**  $\pi_{opt}$  in state  $s$  is the one which gives the largest value for  $Q_{opt}(s)$

$$\pi_{opt}(s) = \arg \max_{a \in Actions(s)} Q_{opt}(s)$$

# Value Iteration

---

$$V_{opt}^{(0)}(s) \leftarrow 0$$

for  $i = 1 \cdots t_{max}$

for each state  $s$

$$V_{opt}^{(t)}(s) \leftarrow \max_{a \in \text{Actions}(s)} \sum_{s'} T(s, a, s') [R(s, \pi(s), s') + \gamma V_{opt}^{(t-1)}(s')]$$

# Value Iteration

---

$V_{opt}^{(0)}(s) \leftarrow 0$   
for  $i = 1 \cdots t_{max}$   
  for each state  $s$   
     $V_{opt}^{(t)}(s) \leftarrow \max_{a \in Actions(s)} \underbrace{\sum_{s'} T(s, a, s') [R(s, \pi(s), s') + \gamma V_{opt}^{(t-1)}(s')]}_{Q_{opt}^{(t-1)}(s)}$

# Value Iteration

---

$V_{opt}^{(0)}(s) \leftarrow 0$   
for  $i = 1 \cdots t_{max}$   
  for each state  $s$   
     $V_{opt}^{(t)}(s) \leftarrow \max_{a \in Actions(s)} \sum_{s'} T(s, a, s') [R(s, \pi(s), s') + \gamma V_{opt}^{(t-1)}(s')]$

**Total states:**  $S$

**Actions per state:**  $A$

**Total successor:**  $S'$

**Complexity:**  $O(SAS't_{max})$

# Value Iteration

---

```
 $V_{opt}^{(0)}(s) \leftarrow 0$   
for  $i = 1 \cdots t_{max}$   
  for each state  $s$   
     $V_{opt}^{(t)}(s) \leftarrow \max_{a \in Actions(s)} \sum_{s'} T(s, a, s') [R(s, \pi(s), s') + \gamma V_{opt}^{(t-1)}(s')]$ 
```

**argmax** instead of **max** will give the optimal policy  $\pi_{opt}$

# Value Iteration

---

Iteration ( $t$ )	$V_{opt}^{(t)}(s_{end})$	$V_{opt}^{(t)}(s_{start})$	$\pi_{opt}(s_{end})$	$\pi_{opt}(s_{start})$
0	0.00	0.00	-	-
1	0.00	10.00	-	Quit
2	0.00	10.67	-	Answer
3	0.00	11.11	-	Answer
100	0.00	12.00	-	Answer

$$V_{\pi}^{(t)}(s_{start}) = 12$$



# Recap

---

$s_{start}$ : start state

**Actions**( $s$ ): all possible actions from state  $s$

**T**( $s, a, s'$ ): probability of  $s'$  if action  $a$  is taken from state  $s$

**Reward**( $s, a, s'$ ): reward from the transition  $s$  to  $s'$

**IsEnd**( $s$ ): is  $s$  a goal state

$0 \leq \gamma \leq 1$ : discount factor (default: 1)

# Unknown Transitions & Reward

---

$s_{start}$ : start state

**Actions**( $s$ ): all possible actions from state  $s$

~~**T**( $s, a, s'$ ): probability of  $s'$  if action  $a$  is taken from state  $s$~~

~~**Reward**( $s, a, s'$ ): reward from the transition  $s$  to  $s'$~~

**IsEnd**( $s$ ): is  $s$  a goal state

$0 \leq \gamma \leq 1$ : discount factor (default: 1)

# Unknown Transitions & Reward

---

$s_{start}$ : start state

**Actions( $s$ )**: all possible actions from state  $s$

~~**T( $s, a, s'$ )**: probability of  $s'$  if action  $a$  is taken from state  $s$~~

~~**Reward( $s, a, s'$ )**: reward from the transition  $s$  to  $s'$~~

**IsEnd( $s$ )**: is  $s$  a goal state

$0 \leq \gamma \leq 1$ : discount factor (default: 1)

## Reinforcement Learning!

# Unknown Transitions & Reward

---

## **MDPs:**

Know how the world works: Environment is observable

Find a policy which maximizes the reward

## **Reinforcement learning:**

Do not know about the world: Environment is not observable

Find a policy which maximizes the reward


Perform actions and collect the reward

# Reinforcement Learning

---

The agent performs actions and observes the rewards

This feedback loop helps learn the missing values (transition probabilities and reward)



# Reinforcement Learning

---

Overall algorithm

for  $t = 1, 2, 3, \dots$

    Choose action  $a_t = \pi_{act}(s_{t-1})$

    Get reward  $r_t$  and new state  $s_t$

    Update parameters

# Model-based Monte Carlo

---

Data:  $s_0; a_1 r_1 s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots$

Estimate  $T(s, a, s')$  &  $R(s, a, s')$


$$\hat{T}(s, a, s') = \frac{\text{No. of times } s, a, s' \text{ occurs}}{\text{No. of times } s, a \text{ occurs}}$$

$$\hat{R}(s, a, s') = \text{reward observed by } s, a, s'$$

# Model-based Monte Carlo

---

Iteration: 0






# Model-based Monte Carlo

Policy  $\pi$  is Answer

Iteration: 1

Data:  $s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$




# Model-based Monte Carlo

Policy  $\pi$  is Answer

Iteration: 2

Data:  $s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$




# Model-based Monte Carlo

Policy  $\pi$  is Answer


Iteration: 3

Data:  $s_{start}$ ;  $Ans, 4, s_{end}$



# Model-based Monte Carlo

---




Can converge to true values

Compute policy using value iteration for the estimated MDP (with  $\hat{T}$  and  $\hat{R}$ )

# Model-based Monte Carlo

---


If  $a \neq \pi(s)$  ( $a = \text{Quit}$ ),  $s, a$  will not be seen



# Model-based Monte Carlo

---

**Exploration:** try unknown actions to get information



# Model-based Monte Carlo

---

We can use the computed transitions and rewards  
And compute the optimal Value and Q-value

$$\hat{V}_{opt}(s) = E[\hat{V}_{opt}(s)] = \begin{cases} 0 & \text{if } isEnd(s) \\ \hat{Q}_{opt}(s) & \text{otherwise} \end{cases}$$

$$\hat{Q}_{opt}(s, a) = \sum_{s'} \hat{T}(s, a, s') [\hat{R}(s, a, s') + \gamma \hat{V}_{opt}(s')]$$

# Model-based Monte Carlo

---

Pros:

- Makes efficient use of experiences

Cons:

- May not scale to large state spaces
  - Learns model one state-action pair at a time
  - Cannot solve MDP for very large  $|S|$



# Model-based vs Model-free

---

Goal: Compute the age of CS students

$P(A)$  is known

$$\begin{aligned}\mathbb{E}[A] &= \sum_a P(A) \cdot a \\ &= 0.35 \times 20 + \dots\end{aligned}$$

# Model-based vs Model-free

---

Without  $P(A)$ , collect samples  $[a_1, a_2, \dots, a_N]$

Unknown  $P(A)$ : *Model-based*

$$\hat{P}(A) = \frac{\text{num}(a)}{N}$$
$$\mathbb{E}[A] \approx \sum_a \hat{P}(A)$$

Because, eventually the correct model is learnt

Unknown  $P(A)$ : *Model-free*

$$\mathbb{E}[A] \approx \frac{1}{N} \sum_i a_i$$

Because, samples appear with right frequencies

# Model-based vs Model-free

---

## **Model based vs. Model free:**

Do we estimate  $T(s, a, s')$  and  $R(s, a, s')$ , or just learn values/policy directly

## **Online vs Batch:**

Learn while exploring the world, or learn from fixed batch of data

## **Active vs Passive:**


Does the learner actively choose actions to gather experience? or, is a fixed policy provided?

# Model-free Monte Carlo

Policy  $\pi$  is Answer

Iteration: 0

Data:




# Model-free Monte Carlo

Policy  $\pi$  is Answer

Iteration: 1

Data:  $s_{start}$ ;  $Ans, 4, s_{end}$




# Model-free Value Iteration

Policy  $\pi$  is Answer

Iteration: 2

Data:  $s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$

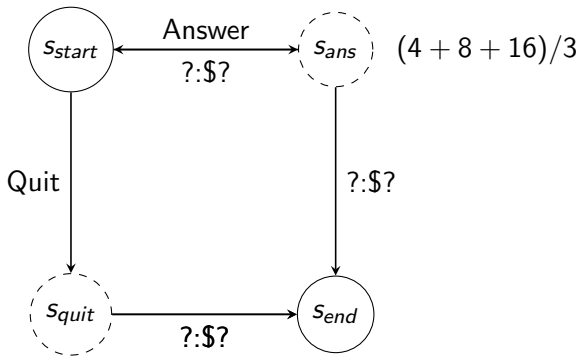


# Model-free Value Iteration

Policy  $\pi$  is Answer

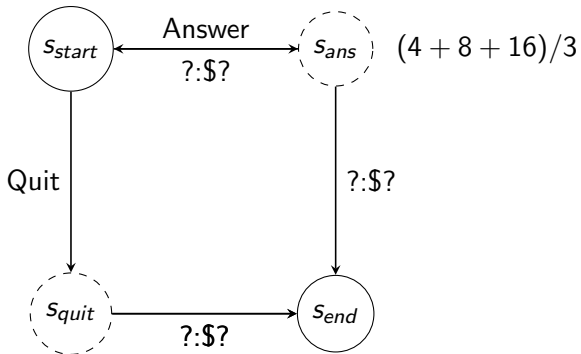
Iteration: 3

Data:  $s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$



# Model-free Value Iteration

We are estimating  $Q_\pi$  and not  $Q_{opt}$





# Model-free Value Iteration

---

**Policy  $\pi$  is Answer**

**Data:**  $s_1; a_1, r_1, s_1; a_2, r_2, s_2; \dots; a_n, r_n, s_n$

$\hat{Q}(s, a) = \text{average of } u_t \text{ where } s_{t-1} = s, a_t = a$

Equivalent formulation (convex combination)

for each  $(s, a, u)$

$$\eta = \frac{1}{1 + \text{No. of updates } (s, a)}$$

$$\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta u$$

# Model-free Value Iteration

---

## Convex combination:

for each  $(s, a, u)$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta u$$

## Stochastic Gradient:

for each  $(s, a, u)$

$$\hat{Q}_{\pi}(s, a) \leftarrow \hat{Q}_{\pi}(s, a) - \eta \left[ \underbrace{\hat{Q}_{\pi}(s, a)}_{\text{prediction}} - \underbrace{u}_{\text{target}} \right]$$

**Objective (Least squares):**  $(\hat{Q}_{\pi}(s, a) - u)^2$

# Using the Utility

---

**Policy  $\pi$  is Answer**

**Data:**

$s_{start}; Ans, 4, s_{end}$	$u = 4$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$u = 8$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$u = 12$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$u = 16$

**Model-free Monte Carlo:**

for each  $(s, a, u)$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta \underbrace{u}_{data}$$

# Using the reward+Q-value

---

Current estimate:  $Q_{\pi}(s, Ans) = 11$

Data:

$s_{start}; Ans, 4, s_{end}$	$4 + 0$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$4 + 11$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$4 + 11$
$s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{start}; Ans, 4, s_{end}$	$4 + 11$

**SARSA:**

for each  $(s, a, r, s', a')$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta[\underbrace{r}_{data} + \gamma \underbrace{\hat{Q}_{\pi}(s', a')}_{estimate}]$$

# Model-free Monte Carlo vs SARSA

---

## Model-free Monte Carlo:

for each  $(s, a, u)$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta \underbrace{u}_{data}$$

## SARSA:

for each  $(s, a, r, s', a')$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta \left[ \underbrace{r}_{data} + \gamma \underbrace{\hat{Q}_{\pi}(s', a')}_{estimate} \right]$$

SARSA uses  $\hat{Q}_{\pi}(s, a)$  instead of raw data  $u$

SARSA doesn't have to wait till it reaches the terminal node to update

# Model-free Monte Carlo vs SARSA

---

Output	MDP	Reinforcement Learning
$Q_\pi$	Policy Evaluation	Model-free Monte Carlo, SARSA
$Q_{opt}$	Value Iteration	<i>Q-Learning</i>

# Q-Learning

---

**Bellman optimality equation:**

$$Q_{opt}(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{opt}(s')]$$

**Q-Learning:**

for each  $(s, a, r, s')$

$$\hat{Q}_{opt}(s, a) \leftarrow (1 - \eta) \underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} + \eta \underbrace{(r + \gamma V_{opt}(s'))}_{\text{target}}$$

# Q-Learning

---

Recall (Bellman optimality equation):

$$Q_{opt}(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{opt}(s')]$$

Q-Learning:

for each  $(s, a, r, s')$

$$\hat{Q}_{opt}(s, a) \leftarrow (1 - \eta) \underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} + \eta \underbrace{(r + \gamma V_{opt}(s'))}_{\text{target}}$$

$$\hat{V}_{opt}(s') = \max_{a' \in \text{Actions}(s')} \hat{Q}_{opt}(s', a')$$



# SARSA vs Q-Learning

---

## SARSA:

for each  $(s, a, r, s', a')$

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta[r + \gamma\hat{Q}_{\pi}(s', a')]$$

## Q-Learning:

for each  $(s, a, r, s')$

$$\hat{Q}_{opt}(s, a) \leftarrow (1 - \eta)\hat{Q}_{opt}(s, a) + \eta(r + \gamma \max_{a' \in \text{Actions}(s')} \hat{Q}_{opt}(s', a'))$$

# Reinforcement Learning

---

**On-policy:** evaluate or improve the data-generating policy

**Off-policy:** evaluate or learn using data from another policy

	On-Policy	Off-Policy
Policy Evaluation ( $Q_{\pi}$ )	Monte-Carlo, SARSA	Q-Learning
Policy Optimization ( $Q_{opt}$ )		

# Reinforcement Learning

---

Algorithm	Estimating	Based On
Model-Based Monte Carlo	$\hat{T}, \hat{R}$	$s_0, a_1, r_1, s_1, \dots$
Model-Free Monte Carlo	$\hat{Q}_\pi$	$u$
SARSA	$\hat{Q}_\pi$	$r + \hat{Q}_\pi$
Q-Learning	$\hat{Q}_{opt}$	$r + \hat{Q}_{opt}$

# Reinforcement Learning

---

Overall algorithm

for  $t = 1, 2, 3, \dots$

    Choose action  $a_t = \pi_{act}(s_{t-1})$

    Get reward  $r_t$  and new state  $s_t$

    Update parameters

# Reinforcement Learning

---

Overall algorithm

for  $t = 1, 2, 3, \dots$

    Choose action  $a_t = \pi_{act}(s_{t-1})$  (**how?**)

    Get reward  $r_t$  and new state  $s_t$

    Update parameters (**how?**)

$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3, \dots; a_n, r_n, s_n$

What policy  $\pi_{act}$  should be used?

# Choosing the policy

---

**Option1:** Select the best policy

$$\pi_{act}(s) = \arg \max_{a \in Actions(s)} \hat{Q}_{\pi}(s, a)$$

**Problem:**  $\hat{Q}_{\pi}(s, a)$  estimates are inaccurate. Too greedy

**Option2:** Select a random policy

$$\pi_{act}(s) = \text{random from } Actions(s)$$

**Problem:** Exploration is not guided

# Epsilon-Greedy Policy

---

$$\pi_{act}(s) = \begin{cases} \arg \max_{a \in Actions(s)} \hat{Q}_{\pi}(s, a) & \text{probability } 1 - \epsilon \\ \text{random from } Actions(s) & \text{probability } \epsilon \end{cases}$$

**A balance between the two!**

# Function Approximation

---

Stochastic Gradient update:

$$\hat{Q}_{opt}(s, a) \leftarrow (1 - \eta)\hat{Q}_{opt}(s, a) + \eta \left[ \underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{opt}(s', a'))}_{\text{target}} \right]$$

How to generalize to unseen states/actions



# Function Approximation

---

## Linear Regression:

Use features  $\phi(s, a)$  and weights  $\mathbf{w}$

$$\hat{Q}_{opt}(s, a; \mathbf{w}) = \mathbf{w} \cdot \phi(s, a)$$

Grid World:

$$\phi_1(s, a) = 1[a = Up]$$

$$\phi_2(s, a) = 1[a = Left]$$

...

$$\phi_7(s, a) = 1[s = (1, *)]$$

$$\phi_8(s, a) = 1[s = (*, 2)]$$

...

# Function Approximation

---

## Q-Learning with Function Approximation:

for each  $(s, a, r, s')$  :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \left[ \underbrace{\hat{Q}_{opt}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{opt}(s'))}_{\text{target}} \right] \phi(s, a)$$

## Objective Function:

$$\left( \underbrace{\hat{Q}_{opt}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{opt}(s'))}_{\text{target}} \right)^2$$

# Recap

---

## **Deterministic vs Stochastic Markov Decision Process**

- Transition
- Reward
- Policy
- Discount

## **Policy value & Q-value Solving MDPs**

- Policy Iteration
- Value Iteration

# Recap

---

**Reinforcement Learning**

**Model-based Monte Carlo Learning**

**Model-free Monte Carlo Learning**

**SARSA**

**Q-Learning**

**Epsilon-Greedy**

**Function Approximation**

# References

---



Stuart Russell and Xiaodong Song (2021)

CS 188 — Introduction to Artificial Intelligence

*University of California, Berkeley*



Chelsea Finn and Nima Anari (2021)

CS221 — Artificial Intelligence: Principles and Techniques

*Stanford University*

# The End